

03_App_Flow

Agentic Regulatory Governance & Unified Simulation

Team Rocket

July 2026

1. High-Level Flow



2. Screen-by-Screen Flow

2.1 Login / Organization Setup

- User signs in (email/password or SSO stub for MVP).
- First-time setup: create organization profile, invite team members, upload initial policy/SOP corpus (bulk upload).

2.2 Dashboard (Home)

- Headline: current **Regulatory Readiness Index** (overall score, gauge visual) + trend line over time.
- Component breakdown: Policy Alignment / Control Coverage / Evidence Completeness / Workflow Readiness / Employee Readiness.
- “High Priority Risks” counter.
- Recent circulars list with per-circular readiness status (Not started / In review / Stress-tested / Remediated).
- Primary CTA: **Upload New Circular**.

2.3 Upload Circular

- Drag-and-drop / file picker for SEBI circular PDF.
- On upload: progress indicator (“Extracting obligations...”).
- System runs OCR fallback automatically if native text extraction is sparse.

2.4 Review Extracted Obligations

- Table view: Obligation description | Deadline | Applicability | Source section (expandable to original text).
- Each row editable/deletable; user can add a missed obligation manually.
- Confirm button locks obligations and proceeds to mapping — this is the human-in-the-loop gate before downstream agents run.

2.5 Policy/Process Mapping Review

- Each confirmed obligation shown with its auto-mapped policy/SOP(s), responsible department, and mapping confidence.
- Obligations with no mapping are visually flagged (“Unmapped — no matching internal control”).
- User can manually re-map or confirm department ownership.

2.6 Run Regulatory Stress Test

- Single CTA: **Run Stress Test**.
- Loading state describes what’s being checked (missing controls, outdated procedures, workflow gaps, missing evidence).
- On completion, routes to Findings view.

2.7 Findings & Gaps View

- List/kanban of findings grouped by severity (High/Medium/Low) and type (Unimplemented / Outdated / Conflict / Workflow gap / Missing evidence).
- Each finding card: short description, affected department, severity badge.
- Click finding → **Regulatory Replay** panel.

2.8 Regulatory Replay (per finding)

- Vertical chain visualization: Circular → Obligation → Mapped Policy → Workflow → Missing Approval/Evidence → Risk → Recommended Fix.
- ARGUS Advisor explanation in plain language beneath the chain.
- CTA: **Add to Action Plan**.

2.9 Evidence Upload / Verification

- Per-obligation or per-finding evidence upload (policy docs, training records, compliance reports, approvals, audit docs).
- Status badges: Present / Stale / Missing, updated as evidence is added.
- Re-running verification updates affected findings and RRI live.

2.10 Action Plan

- Auto-generated prioritized list: Obligation/Gap | Owner (department) | Suggested deadline | Priority.
- Editable — user can reassign owner or adjust deadline.
- Status tracking (Not started / In progress / Done) feeds back into RRI on completion.


2.11 Export Report

- One-click generation of an audit-ready PDF: RRI summary, findings list, Regulatory Replay chains, action plan, evidence status.
- Download / share link.

2.12 ARGUS Advisor (persistent chat)

- Available as a side panel throughout the app.
- Scoped Q&A grounded in the org’s own circulars, policies, and findings (e.g., “Why is our Evidence Completeness only 78?”).

3. Navigation Structure



4. Key Interaction Principles

- **Human confirmation gates** at obligation extraction and mapping steps — AI proposes, human approves, before the Stress Test runs on that data.
- **Everything traceable** — every score, finding, and recommendation links back to source text (no black-box outputs).
- **Progressive disclosure** — dashboard shows the single RRI number first; component scores, findings, and replay chains are drill-downs.
- **Continuous, not one-time** — RRI updates live as evidence is added or action items are completed, reinforcing “continuous readiness monitoring” rather than a static audit snapshot.