

04_UIUX_Brief

Agentic Regulatory Governance & Unified Simulation

Team Rocket

July 2026

1. Design Principles

- Score-first, detail-on-demand.** The Regulatory Readiness Index is the emotional anchor of the product — always visible, always the largest element on the dashboard. Everything else is a drill-down.
- Trust through traceability.** Every AI-generated claim (obligation, gap, score) must visually link back to its source. No unexplained numbers.
- Calm authority, not alarm.** This is a compliance/risk tool — avoid aggressive red/siren aesthetics; use a confident, institutional palette that reads as “credible regulatory technology,” not a security-breach dashboard.
- Reviewable AI, not autonomous AI.** Every AI output has a visible edit/confirm affordance. The UI should never suggest the system acted without human sign-off.

2. Visual Identity

- Palette:**
 - Primary: Deep navy #0F1F3D (institutional trust, regulatory gravitas)
 - Secondary: Slate blue #3A5A8C (interactive elements, links)
 - Accent: Amber #E3A008 (attention/priority — used sparingly for High severity only)
 - Success/Readiness: Teal-green #1F9E85
 - Risk: Muted terracotta #C0503A (not pure red — avoids alarm fatigue)
 - Neutral background: Off-white #F7F8FA, card surface #FFFFFF
- Typography:** A humanist sans-serif for UI text (e.g., Inter) paired with a slightly more geometric/serif display face for the RRI score and report headers to signal “official document” gravitas.
- Iconography:** Line icons, consistent 1.5px stroke — shield/checkmark motifs for readiness, chain-link motif for Regulatory Replay.

3. Core UI Components

3.1 RRI Gauge

- Circular gauge, 0–100, color-graded (terracotta → amber → teal as score rises).
- Center: large numeral + “Regulatory Readiness Index” label.
- Below: 5 component mini-bars (Policy Alignment, Control Coverage, Evidence Completeness, Workflow Readiness, Employee Readiness) with their weights shown on hover.

3.2 Obligation Review Table

- Dense, spreadsheet-like table for compliance-officer efficiency (this persona processes many rows).
- Inline edit, expandable row for source-text citation.
- Bulk-confirm action.

3.3 Finding Card

- Severity badge (color-coded per palette above), finding type tag, one-line description, affected department

chip.

- Click expands into Regulatory Replay.

3.4 Regulatory Replay Chain

- Vertical stepper/timeline component — each node is a distinct entity type (Circular, Obligation, Policy, Workflow, Gap, Fix), connected by directional lines.
- This is the product's signature visual — should feel distinct and ownable, not a generic flowchart.

3.5 Action Plan Board

- Table or lightweight kanban (Not started / In progress / Done) — table preferred for MVP simplicity and for owner/deadline scanning.

3.6 ARGUS Advisor Panel

- Persistent collapsible right-side chat panel, available from any screen, scoped to current org/circular context.

4. Layout & Responsiveness

- **Primary target:** desktop/laptop (compliance officers work in wide, data-dense interfaces).
- Minimum supported width: 1280px for full dashboard; tablet breakpoint (768px) collapses side panels into overlays.
- Mobile is out of scope for MVP beyond basic read-only dashboard viewing.

5. Accessibility

- Color is never the sole indicator of severity/status — always paired with text label or icon (colorblind-safe).
- Minimum WCAG AA contrast for all text on navy/off-white surfaces.
- Keyboard-navigable tables and forms for obligation review (compliance officers will process many rows quickly).

6. Tone of Voice (UI Copy)

- Precise, non-alarmist, non-legalistic. Example: “2 obligations have no matching internal control” rather than “CRITICAL COMPLIANCE FAILURE.”
- Always pair a flagged issue with a next action (“Review mapping” / “Upload evidence”), never leave a gap as a dead-end statement.

7. Demo-Specific UX Notes (Hackathon)

- The 5-minute demo flow (upload → extraction → mapping → stress test → RRI → replay → action plan → report) should be achievable with zero dead clicks — every screen's primary CTA should be the obvious next step in that flow.
- Pre-seed a “before” state (moderate RRI, a few visible gaps) so the live demo shows a believable, non-trivial readiness picture rather than a perfect or empty score.